

Fine-grained ILU Factorization

Group J

Gursimran Singh Saluja

Abdullah Mujahid

Basanagouda Somanakatti

Jun. -Prof. Dr. Andreas Vogel

High Performance Computing in the Engineering Sciences

Computational Engineering

Ruhr Universität Bochum

FINE-GRAINED PARALLEL INCOMPLETE LU FACTORIZATION

EDMOND CHOW AND AFTAB PATEL

*School of Computational Science and Engineering, College of
Computing, Georgia Institute of Technology, Atlanta*

Overview

Introduction : Recap of ILU

Previous Parallel ILU

New Parallel ILU algorithm

- Reformulation of ILU

- Solution of constraint equations

- Algorithms and Implementation

Convergence theory

Experimental Results

- Convergence of the algorithm

- Nonsymmetric, nondiagonally dominant problems

- Results for general SPD problems

- Variation of convergence with problem size

Conventional ILU

- ▶ Given sparse matrix: A
- ▶ Sparsity pattern S :

$$S(i, j) = 1, \text{ for } A(i, j) \neq 0$$

$$S(i, j) = 0, \text{ for } A(i, j) = 0$$

- ▶ Get $A = LU + r$, where L and U also have the same sparsity pattern as A

Conventional ILU

- ▶ Method: Gaussian Elimination
- ▶ Inplace algorithm. L has ones on the diagonal (omitted, only entries of U are stored)
- ▶ Sequential in nature, parallelization is difficult.
- ▶ ILU is mostly used as a *preconditioner*
- ▶ **Important property of ILU :**

$$(LU)_{ij} = A_{ij} \quad \forall (i,j) \in S \quad (1)$$

Previous Parallel ILU

Parallel Strategies

- ▶ Regular ILU isn't a suitable *Preconditioner* for very large matrices
- ▶ Multilevel domain decomposition is preferred
- ▶ ILU implemented in each sub domain (smaller matrices)
- ▶ Parallelism on a single node (on each sub-domain)
- ▶ *Coarse-Grained*

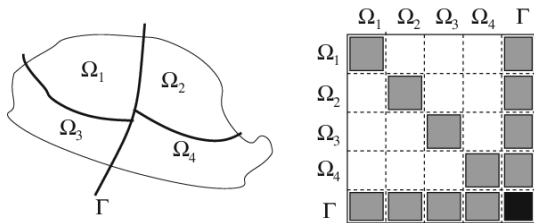


Figure: Domain Decomposition Illustration [3]

New Parallel ILU

Idea

- ▶ New *fine-grained* parallel algorithm implements ILU as solving nonlinear equations
- ▶ Variables : l_{ij}, u_{ij} which are entries of L and U
- ▶ Equations are *constraints* given as follows
- ▶ Entry of ILU is exact on the *Sparsity* pattern S

$$(LU)_{ij} = A_{ij} \quad (2)$$

Problem

- ▶ Variables:



$$l_{ij}, i > j, (i, j) \in S$$



$$u_{ij}, i \leq j, (i, j) \in S$$

- ▶ Normalization: L has unit diagonal (not computed)

- ▶ Number of unknowns: $|S|$



$$\sum_{k=1}^{\min(i,j)} l_{ik} u_{kj} = a_{ij}, (i, j) \in S \quad (3)$$

- ▶ Number of constraints: $|S| = m$ (say)

- ▶ Problem of solving $|S|$ unknowns with $|S|$ equations

- ▶ $|S| > n$, where A is of size $n \times n$

Advantages

- ▶ Equations can be solved in parallel with fine-grained parallelism
- ▶ Exact solution not necessary for *good ILU preconditioner*
- ▶ Good initial guess helps in faster convergence

Solution of constraint equations

- ▶ Explicit expression for each unknown in terms of other unknowns



$$l_{ij} = \frac{1}{u_{ij}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), \quad i > j \quad (4)$$



$$u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, \quad i \leq j \quad (5)$$

- ▶ Form:

$$x^{(p+1)} = G(x^{(p)}), \quad p = 0, 1, \dots$$

- ▶ Initial guess: $x^{(0)}$
- ▶ Components of $x^{(p+1)}$ can be computed in Parallel.

Ordering

$$g(i,j) = \{(i,j) \longrightarrow k\}$$

Row wise ordering

$$\begin{bmatrix} a_{11} & 0 & 0 \\ 0 & a_{22} & 0 \\ a_{31} & 0 & a_{33} \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 3 & 0 & 4 \end{bmatrix}$$

$$g(i,j) = \begin{bmatrix} (1,1) \rightarrow 1 \\ (2,2) \rightarrow 2 \\ (3,1) \rightarrow 3 \\ (3,3) \rightarrow 4 \end{bmatrix}$$

Gaussian Elimination ordering

$$\begin{bmatrix} a_{11} & 0 & 0 \\ 0 & a_{22} & 0 \\ a_{31} & 0 & a_{33} \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 0 & 0 \\ 0 & 3 & 0 \\ 2 & 0 & 4 \end{bmatrix}$$

$$g(i,j) = \begin{bmatrix} (1,1) \rightarrow 1 \\ (2,2) \rightarrow 3 \\ (3,1) \rightarrow 2 \\ (3,3) \rightarrow 4 \end{bmatrix}$$

Gaussian Elimination ordering :

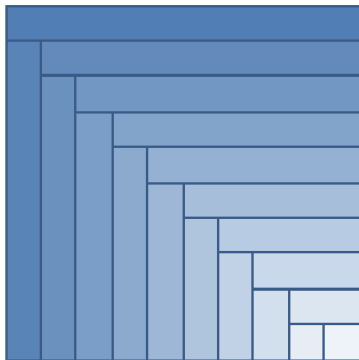


Figure: Ordering illustration

Comparison to exact solution from Gaussian Elimination

Choices

- ▶ Different ordering of the unknown variables l_{ij} , u_{ij} can be chosen
- ▶ Iterative schemes possible:
 - ▶ Asynchronous (parallelizable, Jacobi like)
 - ▶ Synchronous (sequential, Gauss-Seidel like)

Case of Exact ILU from GE

Ordering : Gaussian
Iteration method : Synchronous → Exact ILU in one iteration

Algorithm 1 : Incomplete Factorization

Algorithm 1 Fine-Grained Parallel Incomplete Factorization

```
1: Set unknowns  $x_{g(i,j)}$  ( $l_{ij}$  and  $u_{ij}$ ) to initial values
2: for  $sweep = 1, 2, \dots$  until convergence do
3:   parallel for  $(i, j) \in S$  do
4:     if  $i > j$  then
5:        $x_{g(i,j)} = (a_{ij} - \sum_{k=1}^{j-1} x_{g(i,k)} x_{g(k,j)}) / x_{g(j,j)}$ 
6:     else
7:        $x_{g(i,j)} = (a_{ij} - \sum_{k=1}^{i-1} x_{g(i,k)} x_{g(k,j)})$ 
8:     end if
9:   end
10: end
```

Algorithm 2: Symmetric Incomplete Factorization

Algorithm 2 Symmetric Fine-Grained Parallel IC

```
1: Set unknowns  $x_{g(i,j)}$  ( $l_{ij}$  and  $u_{ij}$ ) to initial values
2: for  $sweep = 1, 2, \dots$  until convergence do
3:   parallel for  $(i, j) \in S_U$  do
4:      $s = a_{ij} - \sum_{k=1}^{j-1} x_{g(i,k)} x_{g(k,j)}$ 
5:     if  $i \neq j$  then
6:        $x_{g(i,j)} = s / x_{g(i,i)}$ 
7:     else
8:        $x_{g(i,i)} = \sqrt{s}$ 
9:     end if
10:  end
11: end
```

Convergence Theory

- ▶ Nonlinear Equation : $F(x) = x - G(x) = 0$
- ▶ Iterative Equation : $x^{(p+1)} = G(x^{(p)})$

Useful Result

Sufficient Condition for **local** linear convergence of fixed point iteration

1. Existence of fixed point (exact ILU guarantees it)
2. G is differentiable around fixed point
3. Spectral radius (maximum eigen value) of $\frac{\partial G}{\partial x}$,

$$\rho\left(\frac{\partial G(x)}{\partial x}\right) < 1$$

Convergence Theory : condition 2 check

$$G_{g(i,j)}(x) = \begin{cases} \frac{1}{x_{g(j,j)}} \left(a_{ij} - \sum_{(i,k), (k,j) \in S} 1 \leq k \leq j-1 x_{g(i,k)} x_{g(k,j)} \right) & \text{if } i > j \\ a_{ij} - \sum_{(i,k), (k,j) \in S} 1 \leq k \leq i-1 x_{g(i,k)} x_{g(k,j)} & \text{if } i \leq j \end{cases} \quad (6)$$

Domain of definition of G :

$$D = \{x \in \mathbb{R} \mid x_{g(j,j)} \neq 0, 1 \leq j \leq n\}$$

Convergence Theory : condition 2 check

For $i > j$:

$$\frac{\partial G_{g(i,j)}}{\partial u_{kj}} = -\frac{l_{ik}}{u_{jj}}, \quad k < j,$$

$$\frac{\partial G_{g(i,j)}}{\partial l_{ik}} = -\frac{u_{kj}}{u_{jj}}, \quad k < j,$$

$$\frac{\partial G_{g(i,j)}}{\partial u_{jj}} = -\frac{1}{u_{jj}^2} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), \quad k < j,$$

Convergence Theory : condition 2 check

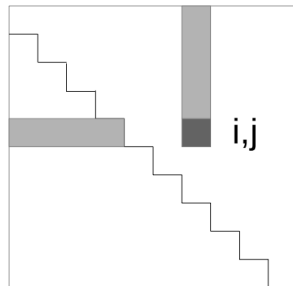
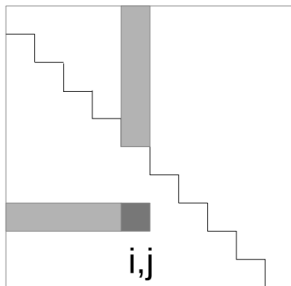
For $i \leq j$:

$$\frac{\partial G_{g(i,j)}}{\partial l_{ik}} = -u_{kj}, \quad k < i,$$

$$\frac{\partial G_{g(i,j)}}{\partial u_{kj}} = -u_{ik}, \quad k < i,$$

Convergence Theory : condition 3 check

Variables dependence :



$$\frac{\partial G(x)}{\partial x} = \begin{bmatrix} 0 & \dots & \dots \\ \vdots & \ddots & \vdots \\ \dots & \dots & 0 \end{bmatrix}_{|S| \times |S|}$$

$$\rho\left(\frac{\partial G(x)}{\partial x}\right) = 0$$

Experimental Results

Test platform

- ▶ Intel Xeon Phi with 61 cores running at 1.09 GHz
- ▶ Each core supporting four way simultaneous multithreading

Aim of tests

1. Convergence of L and U
2. Challenging cases: Nondiagonally dominant
3. Sweeps required for convergence for an effective preconditioner
4. Problem size study
5. Execution times

Experimental Results

Diagnostic tools

- ▶ Convergence is determined by taking 1-norm of the residual

- ▶ L_1 **norm** :

$$\sum_{(i,j) \in S} \left| a_{ij} - \sum_{k=1}^{\min(i,j)} l_{ik} u_{kj} \right|$$

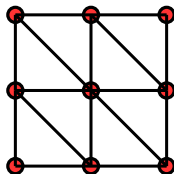
- ▶ **sweep** is the iteration of fixed-point iteration of ILU
- ▶ **solver iteration count** : number of iterations in the PCG algorithm to get the solution of $Ax = b$.
- ▶ Quality of factorization is determined by taking the **solver iteration count**.

Experiment 1 : Convergence of Algorithm

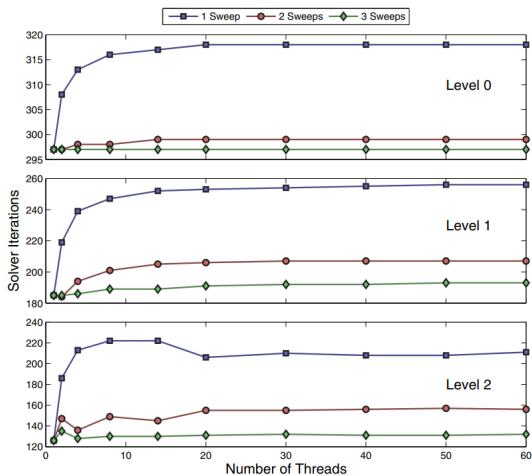
Symmetric Positive Definite Matrix

- ▶ Test Matrix: FEM discretization of Laplacian
- ▶ Matrix entries : 203,841 rows and 1,407,811 nonzeros.
- ▶ Ordering used : Reverse Cuthill-McKee. (Diagonally dominant)
- ▶ Components of b are uniformly distributed from $[-0.5, 0.5]$
- ▶

Sparsity level	0	1	2
Number of non-zeros	805,826	1,008,929	1,402,741

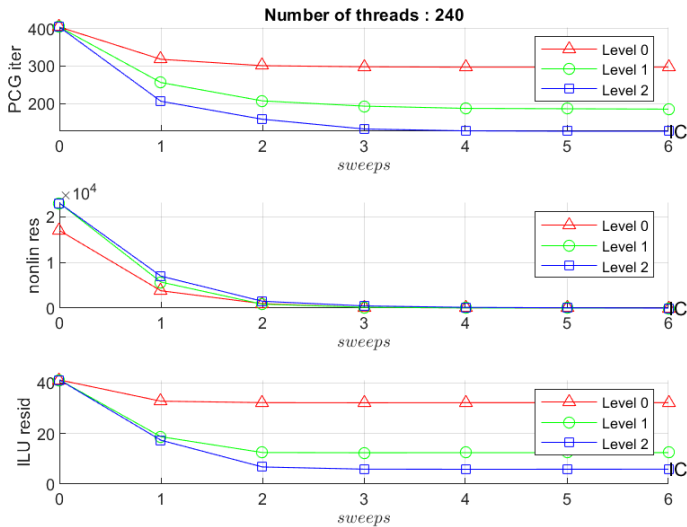


Experiment 1 : Convergence of the algorithm



- ▶ Unpreconditioned case : 1223 iterations
- ▶ No sweeps : 404 iterations

Experiment 1 : Convergence of the algorithm



Experiment 1 : Convergence of the algorithm

Observations

- ▶ Higher level factorizations lead to lower solver iteration counts

level	time for 1 thread	speedup for 60 threads
0	0.189	42.4
1	0.257	44.7
2	0.410	48.8

- ▶ **solver count increase** till 20 threads, then plateau.
- ▶ Algorithm can be highly parallelized
- ▶ Good **preconditioner** is generated with just a single sweep

Experiment 2 : Nonsymmetric, nondiagonally dominant problems

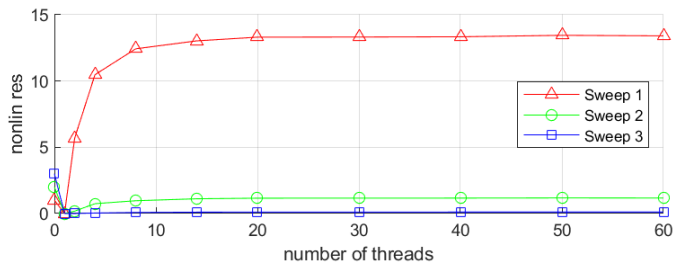
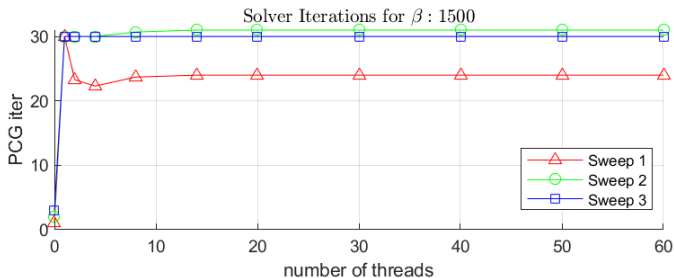
Test Setup

- ▶ 2D Convection-Diffusion problem :

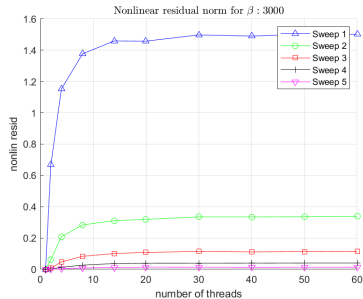
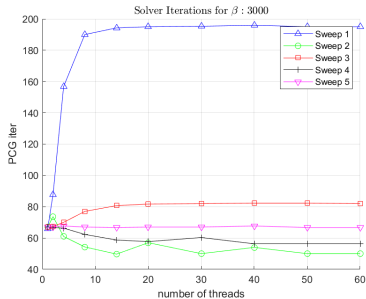
$$-\left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2}\right) + \beta\left(\frac{\partial e^{xy} u}{\partial x} + \frac{\partial e^{-xy} u}{\partial y}\right) = g$$

- ▶ Central Differencing Scheme on $[0, 1] \times [0, 1]$, Dirichlet Boundary Conditions
- ▶ Larger $\beta \rightarrow$ nonsymmetric, non diagonally dominant matrices
- ▶ Mesh Size: 450×450 , Matrix: 202,500 rows and 1,010,700 nonzeros
- ▶ Large $\beta \rightarrow$ more challenging matrices
- ▶ Two tests: $\beta = 1500$ and $\beta = 3000$

Experiment 2 : Nonsymmetric, nondiagonally dominant



Experiment 2 : Nonsymmetric, nondiagonally dominant



Experiment 2 : Nonsymmetric, nondiagonally dominant

Observations

- ▶ No preconditioning CG iterations: 1211 and 1301
- ▶ Nonlinear residuals:
 - ▶ decrease with increasing sweeps
 - ▶ larger for higher thread count
- ▶ Comparable to that from exact IC
- ▶ Few sweeps lead to a good preconditioner

Observations for $\beta = 3000$

- ▶ slow convergence, single sweep no longer is good
- ▶ effect of number of threads is pronounced
- ▶ Algorithm affected by degree of diagonal dominance
- ▶ initial guesses are unstable

Experiment 3 : General SPD Problem

Test Setting

- ▶ Matrices from University of Florida Sparse Matrix Collection
- ▶ level 0 sparsity pattern is used

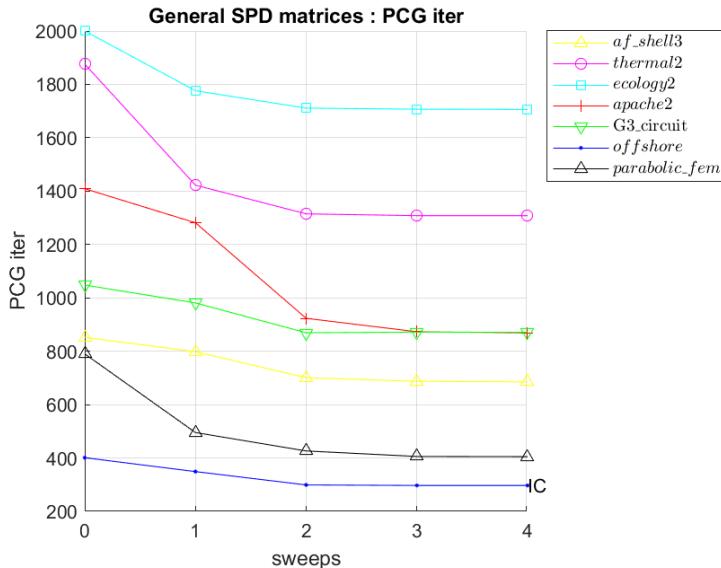
Matrix	No. equations	No. nonzeros
af_shell3	504855	17562051
thermal2	1228045	8580313
ecology2	999999	4995991
apache2	715176	4817870
G3_circuit	1585478	7660826
offshore	259789	4242673
parabolic_fem	525825	3674625

Experiment 3 : General SPD Problem

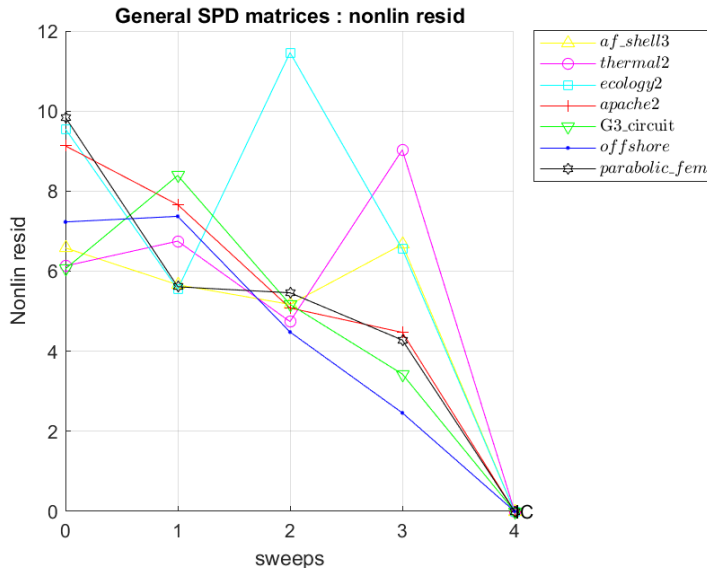
	Sweeps	Nonlin. resid.	PCG iter
af_shell3	0	1.58+05	852
	1	1.66+04	798.3
	2	2.17+03	701
	3	4.67+02	687.3
	IC	0	685
thermal2	0	1.13+05	1876
	1	2.75+04	1422.3
	2	1.74+03	1314.7
	3	8.03+01	1308
	IC	0	1308
ecology2	0	5.55+04	2000+
	1	1.55+04	1776.3
	2	9.46+02	1711
	3	5.55+01	1707
	IC	0	1706
apache2	0	5.13+04	1409
	1	3.66+04	1281.3
	2	1.08+04	923.3
	3	1.47+03	873
	IC	0	869

	Sweeps	Nonlin. resid.	PCG iter
G3_circuit	0	1.06+05	1048
	1	4.39+04	981
	2	2.17+03	869.3
	3	1.43+02	871.7
	IC	0	871
offshore	0	3.23+04	401
	1	4.37+03	349
	2	2.48+02	299
	3	1.46+01	297
	IC	0	297
parabolic_fem	0	5.84+04	790
	1	1.61+04	495.3
	2	2.46+03	426.3
	3	2.28+02	405.7
	IC	0	405

Experiment 3 : General SPD Problem



Experiment 3 : General SPD Problem



Experiment 3 : General SPD Problem

Observations

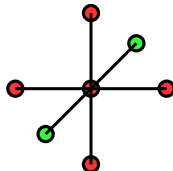
In all cases :

- ▶ solver iteration counts corresponding to : 3 sweeps and exact IC are similar
- ▶ a good preconditioner achieved without full convergence of ILU

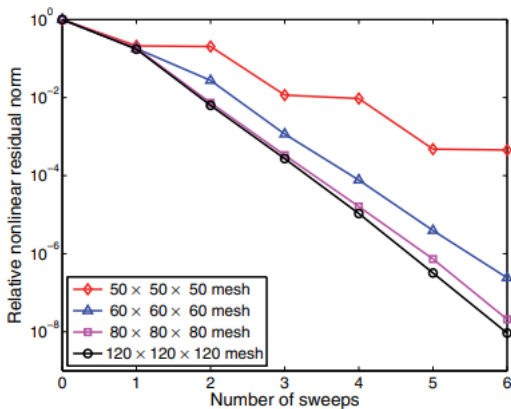
Experiment 4 : Variation of convergence with problem size

Test Setup

- ▶ 7-Point Finite Difference 3D Laplacian matrices
- ▶ Mesh size :
 - ▶ $50 \times 50 \times 50$
 - ▶ $60 \times 60 \times 60$
 - ▶ $70 \times 70 \times 70$
 - ▶ $80 \times 80 \times 80$



Experiment 4 : Convergence with problem size



Experiment 4 : Convergence with problem size

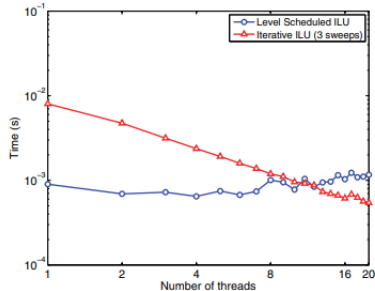
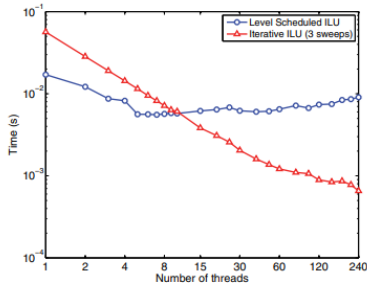
Observations

- ▶ Convergence is better for larger problem size
- ▶ For smaller problems, higher fraction of unknowns updated simultaneously.
- ▶ Asynchronous method being closer to Jacobi type fixed point method.
- ▶ For large problems little variation with problem size

Experiment 5 : Execution time comparison

Test Setup

- ▶ Level scheduled ILU vs Parallel ILU
- ▶ Test case: Matrix from 5-point Finite Difference on 100×100 grid
- ▶ Non-symmetric version algorithm used, ordering is natural



Experiment 5 : Execution time comparison

Result

- ▶ Level scheduled ILU scales well for large parallelism
- ▶ New ILU Algorithm is better when parallelism is limited
- ▶ For large problems with more parallelism, level scheduling is better because parallel ILU uses more sweeps
- ▶ Time for constructing level scheduling is excluded

Approximate Triangular Solves

Parallel Sparse Triangular Solve

- ▶ Time for sparse triangular dominates overall solve time
- ▶ Methods:
 - ▶ Level Schedule Method
 - ▶ Inverse as product of sparse triangular factors



Edmond Chow and Aftab Patel

Fine-Grained Parallel Incomplete LU Factorization,
SIAM J. SCI. COMPUT. Vol. 37, No. 2, pp. C169 - C193,
2015.



Rudi Helfenstein and Jonas Koko

Parallel preconditioned conjugate gradient algorithm on
GPU,
Journal of Computational and Applied Mathematics Volume
236, Issue 15, September 2012, Pages 3584-3590.



Grasedyck, Lars and Kriemann, Ronald and Le Borne,
Sabine

Domain decomposition based \mathcal{H} -LU preconditioning,
Volume 112, Numerische Mathematik, May 2009.



Dr. Vasile Gradinaru and Dr. Roger Käppeli

Numerical methods D-PHYS Course at ETH Zürich
[Course link](#)